

```
In [1]: import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import numpy as np
```

```
In [2]: # ----- 1. 配置实验参数 -----
device = torch.device("cuda" if torch.cuda.is_available() else "cpu") # 优先使用GPU
batch_size = 64
learning_rate = 0.001
epochs = 10 # 训练轮次 (CIFAR10简单, 10轮足够看到效果)
```

```
In [3]: # ----- 2. 定义AlexNet模型 (适配CIFAR10) -----
# 原始AlexNet针对224×224图像, CIFAR10是32×32, 需调整卷积核和池化层参数
class AlexNet(nn.Module):
    def __init__(self, num_classes=10): # CIFAR10有10个类别
        super(AlexNet, self).__init__()
        # 特征提取层: 5个卷积层 + 3个最大池化层
        self.features = nn.Sequential(
            # 卷积层1: 3→64, 核11→3 (适配小图像), 步长4→1, 填充2
            nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=2),
            nn.ReLU(inplace=True),
            # LRN层1: 论文配置n=5, k=2, alpha=1e-4, beta=0.75
            nn.LocalResponseNorm(size=5, alpha=1e-4, beta=0.75, k=2),
            nn.MaxPool2d(kernel_size=2, stride=2), # 输出: 64×16×16

            # 卷积层2: 64 → 192, 核5 → 3, 填充2
            nn.Conv2d(64, 192, kernel_size=3, padding=2),
            nn.ReLU(inplace=True),
            # LRN层2
            nn.LocalResponseNorm(size=5, alpha=1e-4, beta=0.75, k=2),
            nn.MaxPool2d(kernel_size=2, stride=2), # 输出: 192×8×8

            # 卷积层3: 192→384, 核3, 填充1
            nn.Conv2d(192, 384, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            # 卷积层4: 384→256, 核3, 填充1
            nn.Conv2d(384, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            # 卷积层5: 256→256, 核3, 填充1
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2), # 输出: 256×4×4
        )

        # 分类层: 3个全连接层
        self.classifier = nn.Sequential(
            nn.Dropout(0.5), # 防止过拟合
            nn.Linear(256 * 4 * 4, 4096), # 输入: 256×4×4=4096
            nn.ReLU(inplace=True),
```

```

        nn.Dropout(0.5),
        nn.Linear(4096, 4096),
        nn.ReLU(inplace=True),

        nn.Linear(4096, num_classes) # 输出10个类别概率
    )

    def forward(self, x):
        x = self.features(x) # 特征提取
        x = x.view(x.size(0), -1) # 展平 (batch_size, 256×4×4)
        x = self.classifier(x) # 分类
        return x

```

In [5]:

```

# ----- 3. 数据准备 (CIFAR10数据集) -----
# 数据预处理: 归一化、随机增强 (训练集)
transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4), # 随机裁剪
    transforms.RandomHorizontalFlip(), # 随机水平翻转
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
])

train_dataset = datasets.CIFAR10(
    root='./datasets', train=True, download=True, transform=transform_train
)
test_dataset = datasets.CIFAR10(
    root='./datasets', train=False, download=True, transform=transform_test
)

# 数据加载器 (批量读取数据)
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

# CIFAR10类别名称 (用于后续可视化)
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer',
           'dog', 'frog', 'horse', 'ship', 'truck']

```

100.0%

In [6]:

```

# ----- 4. 初始化模型、损失函数、优化器 -----
model = AlexNet(num_classes=10).to(device) # 模型移到GPU/CPU
criterion = nn.CrossEntropyLoss() # 交叉熵损失 (分类任务常用)
optimizer = optim.Adam(model.parameters(), lr=learning_rate) # Adam优化器

# ----- 5. 训练函数 -----
def train(model, train_loader, criterion, optimizer, epoch):
    model.train() # 训练模式 (启用Dropout)
    running_loss = 0.0
    for batch_idx, (inputs, targets) in enumerate(train_loader):
        inputs, targets = inputs.to(device), targets.to(device)

```

```

# 前向传播
outputs = model(inputs)
loss = criterion(outputs, targets)

# 反向传播+参数更新
optimizer.zero_grad() # 清空梯度
loss.backward() # 计算梯度
optimizer.step() # 更新参数

running_loss += loss.item()

# 打印训练进度
if batch_idx % 100 == 99: # 每100个batch打印一次
    print(f'Epoch [{epoch+1}/{epochs}], Batch [{batch_idx+1}/{len(tr
running_loss = 0.0

```

```

In [ ]: # ----- 6. 测试函数 -----
def test(model, test_loader, criterion):
    model.eval() # 测试模式 (禁用Dropout)
    test_loss = 0.0
    correct = 0
    total = 0

    with torch.no_grad(): # 禁用梯度计算 (节省内存)
        for inputs, targets in test_loader:
            inputs, targets = inputs.to(device), targets.to(device)
            outputs = model(inputs)
            loss = criterion(outputs, targets)

            test_loss += loss.item()
            _, predicted = outputs.max(1) # 取概率最大的类别
            total += targets.size(0)
            correct += predicted.eq(targets).sum().item() # 统计正确个数

    # 计算测试集损失和准确率
    avg_loss = test_loss / len(test_loader)
    accuracy = 100. * correct / total
    print(f'\nTest Set: Average Loss: {avg_loss:.4f}, Accuracy: {correct}/{t
    return accuracy

# ----- 7. 启动训练和测试 -----
print(f"Training on {device}...")
best_accuracy = 0.0
# for epoch in range(epochs):
#     train(model, train_loader, criterion, optimizer, epoch)
#     current_accuracy = test(model, test_loader, criterion)

#     # 保存最优模型
#     if current_accuracy > best_accuracy:
#         best_accuracy = current_accuracy
#         torch.save(model.state_dict(), 'results/alexnet_cifar10_best.pth')
#         print(f"Saved best model with accuracy: {best_accuracy:.2f}%\n")

# print(f"Training Finished! Best Accuracy: {best_accuracy:.2f}%")

```

Training on cpu...

```
Epoch [1/10], Batch [100/782], Loss: 2.1783
Epoch [1/10], Batch [200/782], Loss: 1.9728
Epoch [1/10], Batch [300/782], Loss: 1.8626
Epoch [1/10], Batch [400/782], Loss: 1.7660
Epoch [1/10], Batch [500/782], Loss: 1.6747
Epoch [1/10], Batch [600/782], Loss: 1.6193
Epoch [1/10], Batch [700/782], Loss: 1.5739
```

Test Set: Average Loss: 1.5522, Accuracy: 4296/10000 (42.96%)

Saved best model with accuracy: 42.96%

```
Epoch [2/10], Batch [100/782], Loss: 1.5317
Epoch [2/10], Batch [200/782], Loss: 1.4714
Epoch [2/10], Batch [300/782], Loss: 1.4165
Epoch [2/10], Batch [400/782], Loss: 1.4126
Epoch [2/10], Batch [500/782], Loss: 1.3548
Epoch [2/10], Batch [600/782], Loss: 1.3543
Epoch [2/10], Batch [700/782], Loss: 1.3433
```

Test Set: Average Loss: 1.1950, Accuracy: 5690/10000 (56.90%)

Saved best model with accuracy: 56.90%

```
Epoch [3/10], Batch [100/782], Loss: 1.2687
Epoch [3/10], Batch [200/782], Loss: 1.2437
Epoch [3/10], Batch [300/782], Loss: 1.2085
Epoch [3/10], Batch [400/782], Loss: 1.2221
Epoch [3/10], Batch [500/782], Loss: 1.1993
Epoch [3/10], Batch [600/782], Loss: 1.1686
Epoch [3/10], Batch [700/782], Loss: 1.1540
```

Test Set: Average Loss: 1.1346, Accuracy: 5916/10000 (59.16%)

Saved best model with accuracy: 59.16%

```
Epoch [4/10], Batch [100/782], Loss: 1.1187
Epoch [4/10], Batch [200/782], Loss: 1.0914
Epoch [4/10], Batch [300/782], Loss: 1.0966
Epoch [4/10], Batch [400/782], Loss: 1.1024
Epoch [4/10], Batch [500/782], Loss: 1.0780
Epoch [4/10], Batch [600/782], Loss: 1.0473
Epoch [4/10], Batch [700/782], Loss: 1.0496
```

Test Set: Average Loss: 0.9714, Accuracy: 6542/10000 (65.42%)

Saved best model with accuracy: 65.42%

```
Epoch [5/10], Batch [100/782], Loss: 1.0126
Epoch [5/10], Batch [200/782], Loss: 1.0070
Epoch [5/10], Batch [300/782], Loss: 1.0295
Epoch [5/10], Batch [400/782], Loss: 0.9779
Epoch [5/10], Batch [500/782], Loss: 1.0153
Epoch [5/10], Batch [600/782], Loss: 0.9828
Epoch [5/10], Batch [700/782], Loss: 0.9814
```

Test Set: Average Loss: 0.8626, Accuracy: 6936/10000 (69.36%)

Saved best model with accuracy: 69.36%

```
Epoch [6/10], Batch [100/782], Loss: 0.9538
Epoch [6/10], Batch [200/782], Loss: 0.9270
Epoch [6/10], Batch [300/782], Loss: 0.9163
Epoch [6/10], Batch [400/782], Loss: 0.9300
Epoch [6/10], Batch [500/782], Loss: 0.9254
Epoch [6/10], Batch [600/782], Loss: 0.9244
Epoch [6/10], Batch [700/782], Loss: 0.9179
```

Test Set: Average Loss: 0.8268, Accuracy: 7071/10000 (70.71%)

Saved best model with accuracy: 70.71%

```
Epoch [7/10], Batch [100/782], Loss: 0.8774
Epoch [7/10], Batch [200/782], Loss: 0.8996
Epoch [7/10], Batch [300/782], Loss: 0.8732
Epoch [7/10], Batch [400/782], Loss: 0.8915
Epoch [7/10], Batch [500/782], Loss: 0.8569
Epoch [7/10], Batch [600/782], Loss: 0.8615
Epoch [7/10], Batch [700/782], Loss: 0.8778
```

Test Set: Average Loss: 0.8019, Accuracy: 7248/10000 (72.48%)

Saved best model with accuracy: 72.48%

```
Epoch [8/10], Batch [100/782], Loss: 0.8443
Epoch [8/10], Batch [200/782], Loss: 0.8351
Epoch [8/10], Batch [300/782], Loss: 0.8468
Epoch [8/10], Batch [400/782], Loss: 0.8354
Epoch [8/10], Batch [500/782], Loss: 0.8241
Epoch [8/10], Batch [600/782], Loss: 0.8234
Epoch [8/10], Batch [700/782], Loss: 0.7899
```

Test Set: Average Loss: 0.8030, Accuracy: 7221/10000 (72.21%)

```
Epoch [9/10], Batch [100/782], Loss: 0.7887
Epoch [9/10], Batch [200/782], Loss: 0.7999
Epoch [9/10], Batch [300/782], Loss: 0.8060
Epoch [9/10], Batch [400/782], Loss: 0.7931
Epoch [9/10], Batch [500/782], Loss: 0.8028
Epoch [9/10], Batch [600/782], Loss: 0.7951
Epoch [9/10], Batch [700/782], Loss: 0.7935
```

Test Set: Average Loss: 0.7150, Accuracy: 7519/10000 (75.19%)

Saved best model with accuracy: 75.19%

```
Epoch [10/10], Batch [100/782], Loss: 0.7467
Epoch [10/10], Batch [200/782], Loss: 0.7944
Epoch [10/10], Batch [300/782], Loss: 0.7535
Epoch [10/10], Batch [400/782], Loss: 0.7797
Epoch [10/10], Batch [500/782], Loss: 0.7720
```

```
Epoch [10/10], Batch [600/782], Loss: 0.7614
Epoch [10/10], Batch [700/782], Loss: 0.7647

Test Set: Average Loss: 0.7327, Accuracy: 7531/10000 (75.31%)

Saved best model with accuracy: 75.31%

Training Finished! Best Accuracy: 75.31%
```

```
In [ ]: # ----- 8. 单图推理演示 (加载最优模型) -----
def infer_single_image(model_path, image_path, transform):
    # 加载模型
    model = AlexNet(num_classes=10).to(device)
    model.load_state_dict(torch.load(model_path))
    model.eval()

    # 预处理图像
    from PIL import Image
    image = Image.open(image_path).convert('RGB')
    image_tensor = transform(image).unsqueeze(0).to(device) # 增加batch维度

    # 推理
    with torch.no_grad():
        output = model(image_tensor)
        _, predicted = output.max(1)
        predicted_class = classes[predicted.item()]

    # 显示图像和结果
    plt.imshow(image)
    plt.title(f'Predicted: {predicted_class}')
    plt.axis('off')
    plt.show()
```

```
In [10]: # 从test_batch中加载图片
import pickle
def load_cifar10_batch(batch_path):
    with open(batch_path, 'rb') as f:
        batch = pickle.load(f, encoding='bytes')
    images = batch[b'data']
    labels = batch[b'labels']
    images = images.reshape(-1, 3, 32, 32).astype(np.uint8)
    return images, labels
```

```
In [20]: images, labels = load_cifar10_batch('./datasets/cifar-10-batches-py/test_batch')
img, label = images[0], labels[0]
# save the image
from PIL import Image
img1_pil = Image.fromarray(np.transpose(img, (1, 2, 0)))
img1_pil.save('test_image.jpg')
print(label)
```

3

```
In [15]: # 示例: 用一张自己的图片测试 (替换为你的图片路径)
# 注意: 图片尺寸建议32x32 (或自动缩放), 类别需在CIFAR10中
infer_single_image('./results/alexnet_cifar10_best.pth', 'test_image.jpg', t
```

Predicted: cat

